

```
In [2]: %matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import torch
from torch import autograd
import torch.nn.functional as F
import time

images = np.load("D:/work/JHUSchoolStuff/machinelearning/project1/cs475_pr
object_data/images.npy")
labels = np.load("D:/work/JHUSchoolStuff/machinelearning/project1/cs475_pr
object_data/labels.npy")
test = np.load("D:/work/JHUSchoolStuff/machinelearning/project1/cs475_proj
ect_data/test_images.npy")
height = images.shape[1]
width = images.shape[2]
size = height * width
images = (images - images.mean()) / images.std()
data = images.reshape(images.shape[0],size)
data = torch.from_numpy(data).float()
labels = torch.from_numpy(labels).float()
test_data = test.reshape(test.shape[0], size)
test_data = (test_data - test_data.mean()) / test_data.std()
train_seqs = data[0:45000,:]
train_labels = labels[0:45000]
val_seqs = data[45000:,:]
val_labels = labels[45000:]
```

```
In [3]: class TwoLayerNN(torch.nn.Module):
    def __init__(self, layer_1):
        super().__init__()
        self.layer_1 = torch.nn.Linear(height * width, layer_1)
        self.layer_2 = torch.nn.Linear(layer_1, 5)
        self.drop = torch.nn.Dropout(p = 0.3)
    def forward(self, x):
        x = self.layer_1(x)
        y = F.relu(x)
        y = self.drop(y)
        z = self.layer_2(y)
        return z
```

```
In [14]: class ThreeLayerNN(torch.nn.Module):
    def __init__(self, layer_1, layer_2):
        super().__init__()
        self.layer_1 = torch.nn.Linear(height * width, layer_1)
        self.layer_2 = torch.nn.Linear(layer_1, layer_2)
        self.layer_3 = torch.nn.Linear(layer_2, 5)
    def forward(self, x):
        x = self.layer_1(x)
        y = F.relu(x)
        z = self.layer_2(y)
        a = F.relu(z)
        b = self.layer_3(a)
```

```
return b
```

```
In [15]: class FourLayerNN(torch.nn.Module):
    def __init__(self, layer_1, layer_2, layer_3):
        super().__init__()
        self.layer_1 = torch.nn.Linear(height * width, layer_1)
        self.layer_2 = torch.nn.Linear(layer_1, layer_2)
        self.layer_3 = torch.nn.Linear(layer_2, layer_3)
        self.layer_4 = torch.nn.Linear(layer_3, 5)
    def forward(self, x):
        x = self.layer_1(x)
        y = F.relu(x)
        z = self.layer_2(y)
        a = F.relu(z)
        b = self.layer_3(a)
        c = F.relu(b)
        d = self.layer_4(c)
        return d
```

```
In [4]: def train(model, optimizer, batch_size):
    # model.train() puts our model in train mode, which can require different
    # behavior than eval mode (for example in the case of dropout).
    model.train()

    # i is a 1-D array with shape [batch_size]
    i = np.random.choice(train_seqs.shape[0], size=batch_size, replace=False)
    i = torch.from_numpy(i).long()
    x = autograd.Variable(train_seqs[i, :])
    y = autograd.Variable(train_labels[i]).long()
    optimizer.zero_grad()
    y_hat_ = model(x)
    loss = F.multi_margin_loss(y_hat_, y) #using multi_margin_loss for last one
    #loss = F.cross_entropy(y_hat_, y)
    loss.backward()
    optimizer.step()
    return loss.data[0]
```

```
In [5]: def approx_train_accuracy(model):
    model.eval()
    i = np.random.choice(train_seqs.shape[0], size=1000, replace=False)
    i = torch.from_numpy(i).long()
    x = autograd.Variable(train_seqs[i, :])
    y = autograd.Variable(train_labels[i]).long()
    y_hat_ = model(x)
    y_hat = np.zeros(1000)
    for i in range(1000):
        y_hat[i] = torch.max(y_hat_[i, :].data, 0)[1][0]
    return accuracy(y_hat, y.data.numpy())
```

```
In [6]: def val_accuracy(model):
    model.eval()
    x = autograd.Variable(val_seqs)
```

```

y = autograd.Variable(val_labels)
y_hat_ = model(x)
y_hat = np.zeros(5000)
for i in range(5000):
    y_hat[i] = torch.max(y_hat_[i,:].data, 0)[1][0]
return accuracy(y_hat, y.data.numpy())

```

```

In [7]: def accuracy(y, y_hat):
        return (y == y_hat).astype(np.float).mean()

```

```

In [8]: def plot(train_accs, val_accs):
        plt.figure(200)
        plt.title('Training Accuracy')
        plt.xlabel('Iteration')
        plt.ylabel('Accuracy')
        plt.plot(train_accs, 'b')
        plt.show()
        plt.figure(300)
        plt.title('Validation Accuracy')
        plt.xlabel('Iteration')
        plt.ylabel('Accuracy')
        plt.plot(val_accs, 'b')
        plt.show()

```

```

In [9]: def runModel(model, batch_size, NUM_OPT_STEPS, optimizer):
        train_accs, val_accs = [], []
        for i in range(NUM_OPT_STEPS):
            train(model, optimizer, batch_size)
            if i % 100 == 0:
                train_accs.append(approx_train_accuracy(model))
                val_accs.append(val_accuracy(model))
                print("%6d %5.2f %5.2f" % (i, train_accs[-1], val_accs[-1]))
        plot(train_accs, val_accs)

```

```

In [22]: three_layer = ThreeLayerNN(200, 100)
optimizer_three_layer = torch.optim.Adam(three_layer.parameters(), lr=0.001)
runModel(three_layer, 32, 10000, optimizer_three_layer)

```

```


    0  0.37  0.36
  100  0.78  0.79
  200  0.78  0.79
  300  0.80  0.80
  400  0.80  0.81
  500  0.85  0.83
  600  0.83  0.83
  700  0.84  0.83
  800  0.86  0.85
  900  0.87  0.84
 1000  0.87  0.84
 1100  0.86  0.85
 1200  0.87  0.85
 1300  0.88  0.86
 1400  0.85  0.86
 1500  0.88  0.86
 1600  0.89  0.86

```

1700	0.88	0.86
1800	0.88	0.86
1900	0.87	0.86
2000	0.89	0.86
2100	0.87	0.86
2200	0.87	0.86
2300	0.89	0.86
2400	0.89	0.87
2500	0.90	0.86
2600	0.91	0.87
2700	0.89	0.86
2800	0.91	0.87
2900	0.91	0.87
3000	0.91	0.87
3100	0.90	0.87
3200	0.91	0.87
3300	0.91	0.87
3400	0.89	0.87
3500	0.93	0.87
3600	0.92	0.88
3700	0.92	0.87
3800	0.91	0.87
3900	0.92	0.87
4000	0.91	0.87
4100	0.90	0.87
4200	0.93	0.88
4300	0.90	0.87
4400	0.92	0.87
4500	0.92	0.87
4600	0.91	0.88
4700	0.93	0.88
4800	0.92	0.87
4900	0.92	0.86
5000	0.92	0.87
5100	0.93	0.87
5200	0.94	0.88
5300	0.93	0.87
5400	0.94	0.87
5500	0.92	0.87
5600	0.92	0.88
5700	0.94	0.88
5800	0.94	0.88
5900	0.95	0.88
6000	0.93	0.88
6100	0.94	0.88
6200	0.93	0.87
6300	0.94	0.87
6400	0.93	0.87
6500	0.95	0.87
6600	0.94	0.87
6700	0.94	0.88
6800	0.94	0.87
6900	0.94	0.87
7000	0.93	0.87
7100	0.95	0.87
7200	0.95	0.87
7300	0.95	0.87

7400	0.95	0.88
7500	0.95	0.87
7600	0.94	0.87
7700	0.95	0.87
7800	0.95	0.87
7900	0.95	0.87
8000	0.94	0.88
8100	0.94	0.87
8200	0.94	0.87
8300	0.93	0.87
8400	0.94	0.87
8500	0.94	0.88
8600	0.95	0.87
8700	0.96	0.88
8800	0.95	0.87
8900	0.96	0.88
9000	0.95	0.87
9100	0.95	0.87
9200	0.94	0.87
9300	0.96	0.87
9400	0.95	0.87
9500	0.95	0.86
9600	0.96	0.86
9700	0.96	0.87
9800	0.94	0.87
9900	0.96	0.87







```
In [23]: four_layer = FourLayerNN(200, 100, 50)
optimizer_four_layer = torch.optim.Adam(four_layer.parameters(), lr=0.001)
runModel(four_layer, 32, 10000, optimizer_four_layer)
```

0	0.29	0.28
100	0.78	0.76
200	0.77	0.79
300	0.80	0.81
400	0.85	0.82
500	0.83	0.82
600	0.82	0.82
700	0.84	0.84
800	0.82	0.83
900	0.83	0.83
1000	0.85	0.84
1100	0.84	0.85
1200	0.85	0.85
1300	0.88	0.85
1400	0.88	0.86
1500	0.87	0.85
1600	0.86	0.85
1700	0.87	0.86
1800	0.89	0.86
1900	0.85	0.85
2000	0.89	0.86
2100	0.88	0.86
2200	0.87	0.86
2300	0.87	0.86
2400	0.89	0.86
2500	0.89	0.86
2600	0.88	0.86
2700	0.89	0.86
2800	0.88	0.86
2900	0.91	0.87
3000	0.89	0.87
3100	0.89	0.86
3200	0.90	0.87
3300	0.91	0.87
3400	0.92	0.87

3500	0.92	0.87
3600	0.90	0.87
3700	0.90	0.88
3800	0.90	0.87
3900	0.94	0.87
4000	0.90	0.87
4100	0.89	0.87
4200	0.90	0.88
4300	0.90	0.87
4400	0.92	0.87
4500	0.90	0.87
4600	0.93	0.87
4700	0.89	0.87
4800	0.92	0.88
4900	0.92	0.87
5000	0.92	0.88
5100	0.92	0.87
5200	0.91	0.87
5300	0.92	0.88
5400	0.92	0.87
5500	0.93	0.87
5600	0.93	0.87
5700	0.93	0.88
5800	0.93	0.87
5900	0.93	0.88
6000	0.94	0.88
6100	0.93	0.88
6200	0.92	0.87
6300	0.93	0.88
6400	0.93	0.87
6500	0.93	0.87
6600	0.94	0.88
6700	0.91	0.87
6800	0.93	0.87
6900	0.94	0.88
7000	0.94	0.87
7100	0.94	0.87
7200	0.93	0.88
7300	0.94	0.88
7400	0.95	0.88
7500	0.92	0.88
7600	0.95	0.88
7700	0.94	0.87
7800	0.95	0.87
7900	0.95	0.88
8000	0.96	0.88
8100	0.94	0.87
8200	0.95	0.87
8300	0.93	0.87
8400	0.94	0.88
8500	0.95	0.88
8600	0.94	0.87
8700	0.95	0.87
8800	0.94	0.88
8900	0.94	0.87
9000	0.95	0.87
9100	0.94	0.87

9200	0.94	0.87
9300	0.95	0.88
9400	0.95	0.88
9500	0.96	0.87
9600	0.95	0.88
9700	0.94	0.87
9800	0.94	0.87
9900	0.95	0.87




```
In [24]: two_layer = TwoLayerNN(100) #using dropout at 0.3
optimizer_two_layer = torch.optim.Adam(two_layer.parameters(), lr=0.001)
runModel(two_layer, 32, 10000, optimizer_two_layer)
```

0	0.29	0.30
100	0.78	0.78
200	0.78	0.79
300	0.80	0.80
400	0.80	0.81
500	0.82	0.81
600	0.80	0.82
700	0.83	0.82
800	0.83	0.82
900	0.84	0.83

1000	0.84	0.83
1100	0.85	0.83
1200	0.85	0.84
1300	0.86	0.84
1400	0.84	0.84
1500	0.84	0.84
1600	0.85	0.84
1700	0.86	0.84
1800	0.84	0.85
1900	0.83	0.85
2000	0.84	0.85
2100	0.86	0.85
2200	0.85	0.85
2300	0.86	0.85
2400	0.88	0.85
2500	0.86	0.85
2600	0.87	0.86
2700	0.87	0.86
2800	0.87	0.85
2900	0.87	0.86
3000	0.88	0.86
3100	0.88	0.86
3200	0.89	0.86
3300	0.88	0.86
3400	0.88	0.86
3500	0.86	0.86
3600	0.88	0.86
3700	0.88	0.86
3800	0.87	0.86
3900	0.88	0.87
4000	0.87	0.86
4100	0.87	0.86
4200	0.91	0.86
4300	0.87	0.86
4400	0.90	0.86
4500	0.89	0.86
4600	0.90	0.86
4700	0.89	0.86
4800	0.89	0.86
4900	0.88	0.86
5000	0.91	0.87
5100	0.88	0.87
5200	0.88	0.86
5300	0.90	0.87
5400	0.89	0.86
5500	0.90	0.86
5600	0.91	0.87
5700	0.89	0.87
5800	0.88	0.87
5900	0.89	0.87
6000	0.89	0.87
6100	0.89	0.86
6200	0.90	0.87
6300	0.89	0.87
6400	0.91	0.87
6500	0.89	0.87
6600	0.90	0.87

6700	0.89	0.86
6800	0.91	0.87
6900	0.89	0.86
7000	0.89	0.87
7100	0.90	0.86
7200	0.90	0.86
7300	0.90	0.87
7400	0.91	0.87
7500	0.91	0.87
7600	0.90	0.87
7700	0.91	0.87
7800	0.90	0.87
7900	0.90	0.87
8000	0.92	0.87
8100	0.91	0.87
8200	0.91	0.87
8300	0.92	0.87
8400	0.89	0.87
8500	0.90	0.87
8600	0.92	0.87
8700	0.89	0.87
8800	0.92	0.87
8900	0.92	0.87
9000	0.91	0.87
9100	0.93	0.87
9200	0.91	0.87
9300	0.92	0.87
9400	0.91	0.86
9500	0.91	0.87
9600	0.91	0.87
9700	0.91	0.86
9800	0.92	0.86
9900	0.91	0.87






```
In [10]: start = time.time()
two_layer = TwoLayerNN(200) #using dropout at 0.3
optimizer_two_layer = torch.optim.Adam(two_layer.parameters(), lr=0.001)
runModel(two_layer, 65, 10000, optimizer_two_layer)#using multi hinge loss
end = time.time()
print(end - start)
```

0	0.44	0.43
100	0.78	0.79
200	0.81	0.81
300	0.82	0.82
400	0.84	0.83
500	0.83	0.84
600	0.85	0.85
700	0.87	0.83
800	0.86	0.85
900	0.86	0.86
1000	0.86	0.86
1100	0.88	0.86
1200	0.88	0.86
1300	0.89	0.85
1400	0.88	0.85
1500	0.90	0.85
1600	0.89	0.86
1700	0.91	0.86
1800	0.89	0.86
1900	0.89	0.86
2000	0.88	0.87
2100	0.89	0.86
2200	0.89	0.87
2300	0.91	0.87
2400	0.89	0.87
2500	0.90	0.86
2600	0.89	0.87
2700	0.89	0.86
2800	0.90	0.87
2900	0.92	0.87
3000	0.90	0.87
3100	0.90	0.87

3200	0.90	0.87
3300	0.90	0.87
3400	0.91	0.87
3500	0.91	0.87
3600	0.91	0.87
3700	0.92	0.87
3800	0.91	0.87
3900	0.91	0.87
4000	0.93	0.87
4100	0.93	0.88
4200	0.91	0.88
4300	0.92	0.87
4400	0.93	0.88
4500	0.92	0.87
4600	0.94	0.87
4700	0.92	0.88
4800	0.93	0.87
4900	0.93	0.87
5000	0.93	0.88
5100	0.93	0.88
5200	0.93	0.88
5300	0.94	0.87
5400	0.94	0.87
5500	0.93	0.88
5600	0.93	0.87
5700	0.94	0.87
5800	0.95	0.87
5900	0.94	0.88
6000	0.94	0.87
6100	0.93	0.87
6200	0.94	0.87
6300	0.92	0.87
6400	0.93	0.87
6500	0.94	0.87
6600	0.93	0.87
6700	0.94	0.87
6800	0.94	0.87
6900	0.94	0.87
7000	0.94	0.87
7100	0.94	0.88
7200	0.94	0.88
7300	0.95	0.87
7400	0.94	0.87
7500	0.95	0.88
7600	0.94	0.87
7700	0.94	0.87
7800	0.95	0.87
7900	0.95	0.88
8000	0.94	0.87
8100	0.95	0.87
8200	0.94	0.87
8300	0.94	0.88
8400	0.95	0.87
8500	0.96	0.87
8600	0.95	0.87
8700	0.96	0.88
8800	0.95	0.87

8900	0.95	0.87
9000	0.95	0.87
9100	0.96	0.88
9200	0.95	0.87
9300	0.96	0.87
9400	0.95	0.88
9500	0.94	0.88
9600	0.95	0.88
9700	0.95	0.87
9800	0.95	0.88
9900	0.95	0.87



80.23156976699829

Best validation accuracy achieved was 88. This was using my two layer neural net with 200 units, learning rate of 0.001 using Adam optimizer, 10k optimization steps, and 65 batch size. This was also run with a dropout before the second layer with probability 0.3 and using multi hinge loss instead of cross entropy loss. The total time was about 76 seconds.